

004. ChatGPT, fc and conv backprop (fixed).py

```

1 import numpy as np # numerical arrays, linear algebra, randoms – our workhorse
2
3 # =====
4 # Forward-pass primitives
5 # =====
6
7 def conv2d(image, kernel):
8     """
9     Naive 2D 'valid' convolution (really cross-correlation – no kernel flip).
10    image: 2D array (H x W)
11    kernel: 2D array (kH x kW)
12    returns: 2D array of shape ((H - kH + 1) x (W - kW + 1))
13    """
14    h, w = image.shape # input spatial dims
15    kh, kw = kernel.shape # kernel spatial dims
16    out_h = h - kh + 1 # 'valid' conv height
17    out_w = w - kw + 1 # 'valid' conv width
18    output = np.zeros((out_h, out_w)) # allocate result buffer
19
20    # slide kernel over every valid top-left position
21    for i in range(out_h): # row loop over output
22        for j in range(out_w): # col loop over output
23            region = image[i:i+kh, j:j+kw] # current receptive field (kH x kW)
24            output[i, j] = np.sum(region * kernel) # elementwise mult + sum
25    return output # (out_h x out_w)
26
27
28 def relu(x):
29     """
30     Elementwise ReLU: max(0, x) applied to every entry.
31     We use np.maximum (vectorized), not np.max (which would reduce to a scalar).
32     """
33     return np.maximum(0, x) # same shape as x
34

```

```
35
36 def max_pooling(x, size=2, stride=2):
37     """
38     2D max pooling.
39     x: 2D array after ReLU (H x W)
40     size: pooling window (size x size)
41     stride: step between windows
42     returns: pooled map of shape:
43         out_h = floor((H - size)/stride) + 1
44         out_w = floor((W - size)/stride) + 1
45     """
46     h, w = x.shape
47     out_h = (h - size) // stride + 1      # generic formula (not assuming size==stride, though default matches)
48     out_w = (w - size) // stride + 1
49     output = np.zeros((out_h, out_w))    # pooled output
50
51     for i in range(out_h):                # iterate vertical positions
52         for j in range(out_w):            # iterate horizontal positions
53             # extract the pooling window using stride for position and size for window extent
54             region = x[i*stride:i*stride+size, j*stride:j*stride+size]
55             output[i, j] = np.max(region) # route forward the maximum
56     return output
57
58
59 def flatten(x):
60     """
61     Flatten a 2D feature map to 1D vector (row-major).
62     """
63     return x.flatten()                    # shape: (H*W,)
64
65
66 def fully_connected(x, weight, bias):
67     """
68     Dense layer:  $y = Wx + b$ 
69     weight: (num_classes x D)
70     bias: (num_classes,)
71     x: (D,)
```

```
72     returns: (num_classes,)
73     """
74     return np.dot(weight, x) + bias
75
76
77 def softmax(x):
78     """
79     Numerically-stable softmax.
80     x: logits (K,)
81     returns: probabilities summing to 1 (K,)
82     """
83     shifted = x - np.max(x)          # stabilize exponentials
84     exps = np.exp(shifted)           # exp per class
85     return exps / np.sum(exps)       # normalize to simplex
86
87
88 def cross_entropy_loss(probs, label):
89     """
90     Negative log-likelihood for a single label.
91     probs: softmax probs (K,)
92     label: integer class id
93     """
94     return -np.log(probs[label] + 1e-10) # 1e-10 to avoid log(0) tantrums
95
96
97 # =====
98 # Backward-pass (autodiff by hand)
99 # =====
100
101 def grad_fully_connected(x, weights, probs, label):
102     """
103     Gradients for FC layer when loss = cross-entropy(softmax(logits), label).
104     For softmax + CE, dL/dlogits = probs with 1 subtracted at the true class.
105     x: input vector to FC (D,)
106     weights: (K x D)
107     probs: softmax output (K,)
108     label: integer
```

```

109     returns:
110         dW: (K x D), db: (K,), dx: (D,)
111     """
112     dlogits = probs.copy()           # do NOT mutate probs in-place
113     dlogits[label] -= 1.0           # softmax-CE gradient magic
114
115     dW = np.outer(dlogits, x)        # each row k: dlogits[k] * x
116     db = dlogits                     # derivative wrt bias is just dlogits
117     dx = np.dot(weights.T, dlogits)  # push gradient back to input vector
118     return dW, db, dx
119
120
121 def unflatten_gradient(flat_grad, shape):
122     """
123     Reshape a flat gradient vector back to the pooled map shape.
124     flat_grad: (H*W,)
125     shape: tuple (H,W) to restore
126     """
127     return flat_grad.reshape(shape)    # exact inverse of flatten()
128
129
130 def grad_max_pool(dpool_out, relu_out, size=2, stride=2):
131     """
132     Backprop through max-pooling.
133     dpool_out: gradient arriving from above, shape (out_h, out_w)
134     relu_out: the original input to pooling (post-ReLU map), shape (H,W)
135     We re-find argmax in each pooling window and send all gradient to that spot.
136     returns: gradient wrt relu_out, shape (H,W)
137     """
138     d_relu = np.zeros_like(relu_out)  # initialize with zeros (only argmax gets gradient)
139     ph, pw = dpool_out.shape          # pooled spatial dims
140
141     for i in range(ph):               # loop pooled rows
142         for j in range(pw):           # loop pooled cols
143             # slice the exact window that produced pooled value
144             region = relu_out[i*stride:i*stride+size, j*stride:j*stride+size]
145             # index of the maximum inside the region (ties go to the first max)

```

```
146         max_pos = np.unravel_index(np.argmax(region), region.shape)
147         # route upstream gradient ONLY to that max position
148         d_relu[i*stride + max_pos[0], j*stride + max_pos[1]] += dpool_out[i, j]
149     return d_relu
150
151
152 def grad_relu(d_after_relu, pre_relu):
153     """
154     Backprop through ReLU.
155     d_after_relu: gradient wrt ReLU output
156     pre_relu: the tensor BEFORE ReLU (to know where it was <= 0)
157     returns: gradient wrt pre-ReLU input (zero where pre_relu <= 0)
158     """
159     d = d_after_relu.copy()           # avoid mutating caller's buffer
160     d[pre_relu <= 0] = 0.0           # gradient blocked where ReLU was off
161     return d
162
163
164 def grad_conv(image, d_conv_out, kernel_shape):
165     """
166     Gradient wrt the kernel for our 'valid' conv.
167     image: input (H x W)
168     d_conv_out: gradient wrt conv output (H-kH+1 x W-kW+1)
169     kernel_shape: (kH, kW)
170     returns: dKernel (kH x kW)
171     NOTE: We don't compute dImage here (not needed to update weights in this toy).
172     """
173     dkernel = np.zeros(kernel_shape) # accumulator for kernel gradient
174     kh, kw = kernel_shape            # kernel dims
175     dh, dw = d_conv_out.shape         # output gradient dims
176
177     for i in range(dh):               # slide over every output location
178         for j in range(dw):
179             region = image[i:i+kh, j:j+kw] # input patch that contributed
180             dkernel += region * d_conv_out[i, j] # linearity lets us sum contributions
181     return dkernel
182
```

```
183
184 # =====
185 # Tiny training demo (1 step)
186 # =====
187
188 # Reproducibility – keep the RNG civilized
189 np.random.seed(42)
190
191 # Fake input and label for the demo
192 image = np.random.rand(28, 28)          # a pretend grayscale image (MNIST-ish)
193 true_label = 3                          # arbitrarily pick class 3 as the target
194
195 # Initialize parameters (small randoms to avoid early saturation)
196 kernel = np.random.randn(3, 3) * 0.01   # single 3x3 conv filter
197 # After conv (28->26) and 2x2 pool stride 2 (26->13), we have 13*13 features
198 fc_in_dim = 13 * 13
199 num_classes = 10
200 fc_weights = np.random.randn(num_classes, fc_in_dim) * 0.01 # FC weight matrix
201 fc_bias = np.zeros(num_classes)          # bias starts at zeros
202
203 learning_rate = 0.01                     # da tiny SGD step
204
205 # ----- FORWARD PASS -----
206 conv_out = conv2d(image, kernel)          # (26 x 26)
207 relu_out = relu(conv_out)                 # (26 x 26), threshold at 0
208 pool_out = max_pooling(relu_out, size=2, stride=2) # (13 x 13)
209 flat = flatten(pool_out)                  # (169,)
210 logits = fully_connected(flat, fc_weights, fc_bias) # (10,)
211 probs = softmax(logits)                   # (10,), sums to 1
212 loss = cross_entropy_loss(probs, true_label) # scalar
213
214 print("Initial prediction:", np.argmax(probs)) # which class had the max prob
215 print("Loss:", float(loss))                  # cast to float for prettier print
216
217 # ----- BACKWARD PASS -----
218 # Gradients through FC
219 dfc_W, dfc_b, d_flat = grad_fully_connected(flat, fc_weights, probs, true_label) # shapes: (10x169), (10,), (169,)
```

```
220
221 # Reshape gradient back to pooled map shape
222 d_pool = unflatten_gradient(d_flat, pool_out.shape) # (13 x 13)
223
224 # Backprop through max-pool (needs the *forward* relu_out to re-find argmax)
225 d_relu_from_pool = grad_max_pool(d_pool, relu_out, size=2, stride=2) # (26 x 26)
226
227 # Backprop through ReLU to get gradient wrt conv_out (pre-ReLU)
228 d_conv_out = grad_relu(d_relu_from_pool, conv_out) # (26 x 26)
229
230 # Gradient wrt kernel from conv layer
231 dkernel = grad_conv(image, d_conv_out, kernel.shape) # (3 x 3)
232
233 # ----- SGD PARAM UPDATE -----
234 fc_weights -= learning_rate * dfc_W # descend on FC weights
235 fc_bias -= learning_rate * dfc_b # descend on FC bias
236 kernel -= learning_rate * dkernel # descend on conv kernel
237
238 # ----- RE-FORWARD (sanity poke) -----
239 conv_out = conv2d(image, kernel) # recompute with updated params
240 relu_out = relu(conv_out)
241 pool_out = max_pooling(relu_out, size=2, stride=2)
242 flat = flatten(pool_out)
243 logits = fully_connected(flat, fc_weights, fc_bias)
244 probs = softmax(logits)
245 loss = cross_entropy_loss(probs, true_label)
246
247 print("\nAfter one update:")
248 print("Prediction:", np.argmax(probs))
249 print("Loss:", float(loss))
250
```